

HP Helion OpenStack 2.0: Ceph

Contents

HP Helion OpenStack® 2.0: Ceph..... 3

HP Helion OpenStack® 2.0: Ceph

What is Ceph?

Ceph is a Software Defined Storage (SDS) system which comprised of a block storage, object storage, and file system. It is designed to deliver different types of storage interfaces to the end user in the same storage platform. Therefore ceph is termed as unified storage platform.

HP Helion OpenStack : Ceph Overview

Based on OpenStack® Ceph, HP Helion OpenStack® 2.0 Ceph Storage Solution provides an unified scaleable and stable storage solution for the management of Helion OpenStack Volume Storage (Cinder persistent Volumes) service. The solution also supports user backup to the Object Storage (Swift) API service writing to the same unified Ceph storage platform.

This release supports Ceph Firefly 0.80.7 version, which runs on the hlinux kernel 3.14.44-1.

This page describes the integration of Ceph Block Storage with HP Helion OpenStack 2.0.

This guide focuses on installation, configuration and integration between HP Helion OpenStack: 2.0 and Ceph Firefly 0.80.7 running on the hlinux kernel 3.14.44-1.

This guide assumes that you are familiar with the concepts of OpenStack and Ceph. The main purpose of this guide is describe the integration of Ceph Block Storage with HP Helion OpenStack 2.0, detail steps to install dependencies, configure HP Helion OpenStack and Ceph Firefly, and provide troubleshooting guidance.

Installation

This section describes the integration of Ceph Block Storage with HP Helion OpenStack 2.0, detailed procedure to install dependencies, and configure HP Helion OpenStack and Ceph Firefly.

1. Login to the Deployer node.
2. Copy `helion/examples/` in the Deployed node.

```
cp -r ~/helion/examples/ ~/helion-input/my_cloud/definition
```

3. List the folder in `~/helion-input/my_cloud/definition`.

The configuration files for editing are available at `~/helion/my_cloud/definition/data`.

4. Edit the configuration files, based on your environment, to implement Ceph servers.
5. Execute the following command to ensure that the additional disks are available on the servers marked for OSD as specified in the `disks_osd.yml`.

```
vi disks_osd.yml
```

The sample file of `disks_osd.yml` is as follows:

```
disk-models:
- name: DISK_SET_OSD
  # two disk node; remainder of disk 1 and all of disk 2 combined in
  single VG
  # VG is used to create three logical vols for /var, /var/log, and /
  var/crash
  device-groups:
  - name: ceph-osd-data-and-journal
    devices:
    - name: /dev/sdb
    consumer:
      name: ceph
```

```

    attrs:
      usage: data
      journal_disk: /dev/sdc
- name: ceph-osd-data-and-shared-journal-set-1
  devices:
    - name: /dev/sdd
  consumer:
    name: ceph
    attrs:
      usage: data
      journal_disk: /dev/sdf
- name: ceph-osd-data-and-shared-journal-set-2
  devices:
    - name: /dev/sde
  consumer:
    name: ceph
    attrs:
      usage: data
      journal_disk: /dev/sdf

```

The above sample file contains three OSD nodes and two journal disk.

The disk model has the following fields:

device-groups	There can be several device groups. This allows different sets of disks to be used for different purposes.
name	This is an arbitrary name for the device group. The name must be unique.
devices	This is a list of devices allocated to the device group. A name field containing <code>/dev/sdb</code> , <code>/dev/sdd</code> , and <code>/dev/sde</code> indicates that the device group is used by Ceph.
consumer	This specifies the service that uses the device group. A name field containing ceph indicates that the device group is used by Ceph.
attrs	
usage	There can be several use of devices for a particular service. In the above sample, <code>usage</code> field contains data which indicates that the device is used for data storage.
journal_disk	what is its usage?? You can share the journal disk between two nodes.



Important: Minimum 3 nodes are required as OSD nodes.

- Commit your configuration to a [local repo](#):

```

cd ~/helion/hos/ansible
git add -A
git commit -m "<commit message>"

```



Note: Enter your commit message `<commit message>`

- Run the configuration processor

```

cd ~/helion/hos/ansible

```

```
ansible-playbook -i hosts/localhost config-processor-run.yml
```

8. Use `ansible-playbook -i hosts/localhost ready-deployment.yml` file to create a deployment directory.
9. Run `verb_host` commands from the following directory:

```
~/scratch/ansible/next/hos/ansible
```

10. Modify `./helion/hlm/ansible/hlm-deploy.yml` to uncomment the line containing `ceph-deploy.yml`.
11. Run the following ansible playbook:

```
ansible-playbook -i hosts/verb_hosts site.yml
```

Ceph Monitor service is deployed on the Controller Nodes and OSD's are deployed as separate nodes (Resource Nodes).

Run Ceph Client Packages

Execute the following command to install the ceph client packages on controller nodes and create users and ceph pools on the resource nodes:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ceph-client-prepare.yml
```

Configure Ceph as a Cinder backend

Perform the following procedure on the Deployer node to configure Ceph as a Cinder backend:

1. Edit `~/helion/hos/ansible/roles/_CND-CMN/templates/cinder.conf.j2` to add ceph configuration data as shown below:

```
enabled_backends=ceph1
```

2. Copy the following configurations:

```
[ceph1]
rbd_max_clone_depth = 5
rbd_flatten_volume_from_snapshot = False
rbd_uuid = 457eb676-33da-42ec-9a8c-9293d545c337
rbd_user = cinder
rbd_pool = volumes
rbd_ceph_conf = /etc/ceph/ceph.conf
volume_driver = cinder.volume.drivers.rbd.RBDDriver
volume_backend_name = ceph
```



Note: The `rbd_uuid` is available in `"/home/stack/helion/hos/ansible/roles/ceph-client-prepare/vars/ceph_user_model.yml"`

3. Modify `cinder.conf.j2` at `~/helion/hos/ansible/roles/_CND-CMN/templates/cinder.conf.j2` with the following values:

```
backup_driver = cinder.backup.drivers.ceph
backup_ceph_conf = /etc/ceph/ceph.conf
backup_ceph_user = cinder-backup
backup_ceph_chunk_size = 134217728
backup_ceph_pool = backups
backup_ceph_stripe_unit = 0
backup_ceph_stripe_count = 0
restore_discard_excess_bytes = true
```

Parameter description required

4. On all the Controller nodes copy the following packages:

```
cp /usr/lib/python2.7/dist-packages/rbd.py /opt/stack/venv/
cinder-20150827T030317Z/lib/python2.7/site-packages/
cp /usr/lib/python2.7/dist-packages/rados.py /opt/stack/venv/
cinder-20150827T030317Z/lib/python2.7/site-packages/
```



Note: Currently RBD volume attach to a nova instance is **NOT** working.

5. Copy `ceph.client.cinder.keyring` to the controller nodes:

- a. Login to controller node as a root user and execute the following command.

```
# ceph auth get-or-create client.cinder | tee /etc/ceph/
ceph.client.cinder.keyring
```

OR

You can execute the following command from the deployer node.

```
# cp /etc/ceph/ceph.client.cinder.keyring to the /etc/ceph folder on the
controller nodes.
```

6. Commit your configuration to the local repo to configure cinder on the deployer node

```
cd /home/stack/helion/hos/ansible
git add -A
git commit -m "<your commit message>"
```



Note: Enter your commit message <commit message>

7. Use `ansible-playbook -i hosts/localhost ready-deployment.yml` file to create a deployment directory.
8. Run `verb_host` commands from the following directory:

```
~/scratch/ansible/next/hos/ansible
```

9. Run the following ansible playbook:

```
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

Once cinder is configured, launch the Horizon dashboard to create a cinder volume type.

Creating Cinder Volume Type

To create a volume type using the Horizon dashboard, do the following:

1. Log into the Horizon dashboard. The Horizon dashboard displays with the options in the left panel.
2. From the left panel, click the **Admin** tab and then click the **Volumes** tab to display the Volumes page.

HP Helion OpenStack® admin

Project

Admin

System

Overview

Hypervisors

Host Aggregates

Instances

Volumes

Flavors

Images

Networks

Volumes

Volumes Volume Types Volume Snapshots

Volume Types

<input type="checkbox"/>	Name	Ass
Displaying 0 items		

QOS Specs

	Name	Consumer
Displaying 0 items		

- Click **Create Volume Type** to display a dialog box.

Create Volume Type

Name *

Description:

Volume type is a type or label that can be selected at volume creation time in OpenStack. It usually maps to a set of capabilities of the storage back-end driver to be used for this volume. Examples: "Performance", "SSD", "Backup", etc. This is equivalent to the `cinder type-create` command. Once the volume type gets created, click the "View Extra Specs" button to set up extra specs key-value pair(s) for that volume type.

Cancel Create Volume Type

- Enter the name of the volume type.
- Click **Create Volume Type**. The newly created volume displays in the Volumes page.

HP Helion OpenStack® admin

Project

Admin

System

Overview

Hypervisors

Host Aggregates

Instances

Volumes

Flavors

Images

Networks

Volumes

Volumes Volume Types Volume Snapshots

Volume Types

<input type="checkbox"/>	Name	Ass
<input type="checkbox"/>	rd	

Displaying 1 items

QOS Specs

	Name	Consumer

Displaying 0 items

Associate the volume type to a backend

To map a volume type to a backend, do the following:

1. Login to the Overcloud Horizon dashboard. The Overcloud dashboard displays with the options in the left panel.

- From the left panel, click the **Admin** tab and then click the **Volumes** tab to display the Volumes

HP Helion OpenStack®

admin

Project

Admin

System

Overview

Hypervisors

Host Aggregates

Instances

Volumes

Flavors

Images

Networks

Volumes

Volume Types

Volume Snapshots

Volume Types

	Name
<input type="checkbox"/>	rbd

Displaying 1 items

QOS Specs

	Name	Consumer

Displaying 0 items

page.

3. Click **View Extra Specs** displayed against the volume type which you want to associate to the

The screenshot shows the HP Helion OpenStack® interface. The left sidebar has a 'Project' dropdown, an 'Admin' dropdown, and a 'System' dropdown. The 'Volumes' link is highlighted in the sidebar. The main content area is titled 'Volumes' and has three tabs: 'Volumes', 'Volume Types', and 'Volume Snapshots'. The 'Volume Types' tab is selected, displaying a table with one row: 'rbd'. Below the table, it says 'Displaying 1 items'. Underneath is a 'QOS Specs' section with a table that has two columns: 'Name' and 'Consumer'. The table is currently empty, and it says 'Displaying 0 items'.

backend.

The **Create Volume Type Extra Specs** dialog box displays.

The screenshot shows the 'Create Volume Type Extra Spec' dialog box. It has a title bar with a close button. The main area has two input fields: 'Key: *' and 'Value: *'. The 'Key' field contains the text 'volume backend name' and the 'Value' field contains 'helion-ceph-cinder'. To the right of these fields is a 'Description:' section with the text: 'Create a new "extra spec" key-value pair for a volume type.' At the bottom right of the dialog are two buttons: 'Cancel' and 'Create'.

4. In the **Key** box, enter *volume_backend_name*. This is the name of the key used to specify the storage backend when provisioning volumes of this volume type.
5. In the **Value** box, enter the name of the backend to which you want to associate the volume type. For example: *helion-ceph-cinder*.
6. Click **Create** to create the extra volume type specs. Note: Once the volume type is mapped to the backend, you can create volumes.



Note: Once the volume type is mapped to the backend, you can create volumes.