

Installing Ceph

Contents

Installing Ceph..... 3

Installing Ceph

This page describes configuration, installation, and the integration of Ceph Block Storage with HP Helion OpenStack 2.0. It also provides the procedure to perform Ceph operation after Ceph deployment.

Prerequisite

The deployer node must be setup before deploying Ceph.

Installation Procedure

Perform the following procedures to configuration, installation, and the integration of Ceph Block Storage with HP Helion OpenStack 2.0

1. Login to the Deployer node.
2. Copy `helion/examples/` in the Deployed node.

```
cp -r ~/helion/examples/ ~/helion-input/my_cloud/definition
```

3. List the folder in `~/helion-input/my_cloud/definition`.

The configuration files for editing are available at `~/helion/my_cloud/definition/data`.

4. Edit the configuration files, based on your environment, to implement Ceph servers.
5. Execute the following command to ensure that the additional disks are available on the servers marked for OSD as specified in the `disks_osd.yml`.

```
vi disks_osd.yml
```

The sample file of `disks_osd.yml` is as follows:

```
disk-models:
- name: DISK_SET_OSD
  # two disk node; remainder of disk 1 and all of disk 2 combined in
  single VG
  # VG is used to create three logical vols for /var, /var/log, and /
  var/crash
  device-groups:
  - name: ceph-osd-data-and-journal
    devices:
    - name: /dev/sdb
    consumer:
      name: ceph
      attrs:
        usage: data
        journal_disk: /dev/sdc
  - name: ceph-osd-data-and-shared-journal-set-1
    devices:
    - name: /dev/sdd
    consumer:
      name: ceph
      attrs:
        usage: data
        journal_disk: /dev/sdf
  - name: ceph-osd-data-and-shared-journal-set-2
    devices:
    - name: /dev/sde
    consumer:
      name: ceph
      attrs:
```

```
usage: data
journal_disk: /dev/sdf
```

The above sample file contains three OSD nodes and two journal disk.


The disk model has the following fields:

| | |
|----------------------|---|
| device-groups | There can be several device groups. This allows different sets of disks to be used for different purposes. |
| name | This is an arbitrary name for the device group. The name must be unique. |
| devices | This is a list of devices allocated to the device group. A name field containing <code>/dev/sdb</code> , <code>/dev/sdd</code> , and <code>/dev/sde</code> indicates that the device group is used by Ceph. |
| consumer | This specifies the service that uses the device group. A name field containing ceph indicates that the device group is used by Ceph. |
| attrs | This is the list of attributes. |
| usage | There can be several use of devices for a particular service. In the above sample, <code>usage</code> field contains data which indicates that the device is used for data storage. |
| journal_disk | It is to used to capture journal data. You can share the journal disk between two nodes. |

 **Important:** Minimum 3 OSD nodes are required to configure Ceph.

- Commit your configuration to a *local repository*:

```
cd ~/helion/hos/ansible
git add -A
git commit -m "<commit message>"
```

 **Note:** Enter your commit message `<commit message>`

- Run the configuration processor

```
cd ~/helion/hos/ansible
ansible-playbook -i hosts/localhost config-processor-run.yml
```

- Use `ansible-playbook -i hosts/localhost ready-deployment.yml` file to create a deployment directory.
- Run `verb_host` commands from the following directory:

```
~/scratch/ansible/next/hos/ansible
```

- Modify `./helion/hlm/ansible/hlm-deploy.yml` to uncomment the line containing `ceph-deploy.yml`.
- Run the following ansible playbook:

```
ansible-playbook -i hosts/verb_hosts site.yml
```

Ceph Monitor service is deployed on the Controller Nodes and OSD's are deployed as separate nodes (Resource Nodes).

Run Ceph Client Packages

Execute the following command to install the ceph client packages on controller nodes and create users and ceph pools on the resource nodes:

```
cd ~/scratch/ansible/next/hos/ansible
ansible-playbook -i hosts/verb_hosts ceph-client-prepare.yml
```

Configure Ceph as a Cinder backend

Perform the following procedure on the Deployer node to configure Ceph as a Cinder backend:

1. Edit `~/helion/hos/ansible/roles/_CND-CMN/templates/cinder.conf.j2` to add ceph configuration data as shown below:

```
enabled_backends=ceph1
```

2. Copy the following configurations:

```
[ceph1]
rbd_max_clone_depth = 5
rbd_flatten_volume_from_snapshot = False
rbd_uuid = 457eb676-33da-42ec-9a8c-9293d545c337
rbd_user = cinder
rbd_pool = volumes
rbd_ceph_conf = /etc/ceph/ceph.conf
volume_driver = cinder.volume.drivers.rbd.RBDDriver
volume_backend_name = ceph
```



Note: The `rbd_uuid` is available in `/home/stack/helion/hos/ansible/roles/ceph-client-prepare/vars/ceph_user_model.yml`

3. Modify `cinder.conf.j2` at `~/helion/hos/ansible/roles/_CND-CMN/templates/cinder.conf.j2` with the following values:

```
backup_driver = cinder.backup.drivers.ceph
backup_ceph_conf = /etc/ceph/ceph.conf
backup_ceph_user = cinder-backup
backup_ceph_chunk_size = 134217728
backup_ceph_pool = backups
backup_ceph_stripe_unit = 0
backup_ceph_stripe_count = 0
restore_discard_excess_bytes = true
```

4. On all the Controller nodes copy the following packages:

```
cp /usr/lib/python2.7/dist-packages/rbd.py /opt/stack/venv/
cinder-20150827T030317Z/lib/python2.7/site-packages/
cp /usr/lib/python2.7/dist-packages/rados.py /opt/stack/venv/
cinder-20150827T030317Z/lib/python2.7/site-packages/
```



Important: Beta1 does not support the RBD volume attachment to a Nova instance.

5. Copy `ceph.client.cinder.keyring` to the controller nodes:
 - a. Login to controller node as a root user and execute the following command.

```
ceph auth get-or-create client.cinder | tee /etc/ceph/
ceph.client.cinder.keyring
```

OR

You can execute the following command from the deployer node.

```
cp /etc/ceph/ceph.client.cinder.keyring to the /etc/ceph folder on the controller nodes.
```

6. Commit your configuration to the local repository to configure cinder on the deployer node

```
cd /home/stack/helion/hos/ansible
git add -A
git commit -m "<your commit message>"
```



Note: Enter your commit message <commit message>

7. Use `ansible-playbook -i hosts/localhost ready-deployment.yml` file to create a deployment directory.
8. Run `verb_host` commands from the following directory:

```
~/scratch/ansible/next/hos/ansible
```

9. Run the following ansible playbook:

```
ansible-playbook -i hosts/verb_hosts cinder-reconfigure.yml
```

Once cinder is configured, launch the Horizon dashboard to create a cinder volume type.

Creating Cinder Volume Type

To create a volume type using the Horizon dashboard, do the following:

1. Log into the Horizon dashboard. The Horizon dashboard displays with the options in the left panel.
2. From the left panel, click the **Admin** tab and then click the **Volumes** tab to display the Volumes page.

The screenshot shows the HP Helion OpenStack® admin interface. The top navigation bar includes the logo and a user menu for 'admin'. The left sidebar contains a navigation menu with categories: Project, Admin, and System. Under the Admin category, the 'Volumes' option is selected. The main content area is titled 'Volumes' and has three tabs: 'Volumes', 'Volume Types', and 'Volume Snapshots'. The 'Volume Types' tab is active, displaying a table with columns 'Name' and 'Ass'. Below the table, it says 'Displaying 0 items'. There is also a section for 'QOS Specs' with a table with columns 'Name' and 'Consumer', also showing 'Displaying 0 items'.

3. Click **Create Volume Type** to display a dialog box.

The screenshot shows a 'Create Volume Type' dialog box. It has a title bar with a close button. The main area contains a 'Name' field with a '*' indicating it is required, and a 'Description' section. The description text reads: 'Volume type is a type or label that can be selected at volume creation time in OpenStack. It usually maps to a set of capabilities of the storage back-end driver to be used for this volume. Examples: "Performance", "SSD", "Backup", etc. This is equivalent to the cinder type-create command. Once the volume type gets created, click the "View Extra Specs" button to set up extra specs key-value pair(s) for that volume type.' At the bottom right, there are two buttons: 'Cancel' and 'Create Volume Type', with the latter highlighted by a red rectangle.

4. Enter the name of the volume type.
5. Click **Create Volume Type**. The newly created volume displays in the Volumes page.

HP Helion OpenStack® admin

Project Admin System

Overview
Hypervisors
Host Aggregates
Instances
Volumes
Flavors
Images
Networks

Volumes

Volumes Volume Types Volume Snapshots

Volume Types

| <input type="checkbox"/> | Name | Ass |
|--------------------------|------|-----|
| <input type="checkbox"/> | rbd | |

Displaying 1 items

QOS Specs

| | Name | Consumer |
|--|------|----------|
| | | |

Displaying 0 items

Associate the volume type to a backend

To map a volume type to a backend, do the following:

1. Login to the Opencloud Horizon dashboard. The Opencloud dashboard displays with the options in the left panel.
2. From the left panel, click the **Admin** tab and then click the **Volumes** tab to display the Volumes page.

The screenshot shows the HP Helion OpenStack Admin interface. The top header includes the logo and a user menu for 'admin'. The left sidebar contains a navigation menu with categories: Project, Admin, and System. Under Admin, the 'Volumes' option is selected. The main content area is titled 'Volumes' and has three tabs: 'Volumes', 'Volume Types', and 'Volume Snapshots'. The 'Volume Types' tab is active, displaying a table with one entry: 'rbd'. Below the table, it says 'Displaying 1 items'. Underneath the 'Volume Types' section is a 'QOS Specs' section with a table that has columns 'Name' and 'Consumer'. It currently shows 'Displaying 0 items'.

HP Helion OpenStack® admin

Project Admin System

Overview
Hypervisors
Host Aggregates
Instances
Volumes
Flavors
Images
Networks

Volumes

Volumes Volume Types Volume Snapshots

Volume Types

| <input type="checkbox"/> | Name | Ass |
|--------------------------|------|-----|
| <input type="checkbox"/> | rbd | |

Displaying 1 items

QOS Specs

| | Name | Consumer |
|--|------|----------|
|--|------|----------|

Displaying 0 items

3. Click **View Extra Specs** displayed against the volume type which you want to associate to the backend.

HP Helion OpenStack® admin

Volumes

Volumes | Volume Types | Volume Snapshots

Volume Types

| <input type="checkbox"/> | Name | Ass |
|--------------------------|------|-----|
| <input type="checkbox"/> | rbd | |

Displaying 1 items

QOS Specs

| | Name | Consumer |
|--|------|----------|
| | | |

Displaying 0 items

The **Create Volume Type Extra Specs** dialog box displays.

Create Volume Type Extra Spec

Key: *

volume backend name

Value: *

helion-ceph-cinder

Description:
Create a new "extra spec" key-value pair for a volume type.

Cancel Create

- In the **Key** box, enter `volume_backend_name`. This is the name of the key used to specify the storage backend when provisioning volumes of this volume type.

5. In the **Value** box, enter the name of the backend to which you want to associate the volume type. For example: *helion-ceph-cinder*.
6. Click **Create** to create the extra volume type specs.



Note: Once the volume type is mapped to the backend, you can create volumes.

Ceph Operations

After the successful deployment of Ceph, you can perform the following ceph operations:

- Check the status of Ceph OSD and Monitor Services
- Start OSD Nodes and Monitor Services
- Stop OSD Nodes and Monitor Services

Check the status of Ceph OSD and Monitor Services

Perform the following steps to check the status of Ceph OSD Nodes and Monitor Services:

1. Login to Deployer Node.
2. Execute the following command:

```
cd ~/scratch/ansible/next/hos/ansible
```

3. Run the ansible playbook:

```
ansible-playbook -i hosts/verb_hosts ceph-status.yml
```

Start OSD and Monitor Services

Perform the following steps to start OSD Nodes and Monitor Services:

1. Login to Deployer Node.
2. Execute the following command:

```
cd ~/scratch/ansible/next/hos/ansible
```

3. Run the ansible playbook:

```
ansible-playbook -i hosts/verb_hosts ceph-start.yml
```

Stop OSD Nodes and Monitor Services

Perform the following steps to stop OSD Nodes and Monitor Services:

1. Login to Deployer Node.
2. Execute the following command:

```
cd ~/scratch/ansible/next/hos/ansible
```

3. Run the ansible playbook:

```
ansible-playbook -i hosts/verb_hosts ceph-stop.yml
```