

Deploy ESX Cloud

Contents

Deploy ESX Cloud..... 3

Deploy ESX Cloud

Known Restrictions

- Ensure the [minimum hardware requirements](#) are met.
- All disks on the system(s) will be wiped: /dev/sda will be used for the OS but all other disks will be wiped.
- The deployer node must use the HP Linux for HP Helion OpenStack ISO, which can be downloaded from the [Helion Downloads](#) page.
- Three NIC are tested, one 1G used for PXE and two bonded 10G for everything else. All machines can net boot from PXE and use the deployer as a DHCP server.
- All machines of a single type should be the same, that is, all computes, and so on.
- The deployer node must be a dedicated node in Beta1.
- The machine hosting the deployer and all baremetal systems must be connected to a management network. Nodes on this management network must be able to reach the iLO subsystem of each baremetal system to enable host reboots as part of the install process. The HP Helion OpenStack architecture requires that the IPMI network is a separate network and that a route exists from the management network to the IPMI network for iLO access.

Before You Start

1. Prepare your baremetal hardware, as follows, on all nodes:
 - Set up the iLO Advanced license in the iLO configuration.
 - Switch from UEFI to Legacy BIOS.
 - Ensure that the network to be used for PXE installation has PXE enabled.
 - Ensure that the other networks have PXE disabled.
 - Insert the CD ROM in the Virtual Media drive on the iLO.
2. ESX/vCenter integration is not fully automatic, vCenter administrators are advised of the following responsibilities to ensure secure operation:
 - The VMware administrator is responsible for administration of the vCenter servers and the ESX nodes using the VMware administration tools. These responsibilities include:
 - Installing and configuring vCenter Server
 - Installing and configuring ESX server and ESX cluster
 - Installing and configuring shared datastores
 - Establishing network connectivity between the ESX network and the HP Helion management network
 - The VMware administration staff is responsible for the review of vCenter logs. These logs are not automatically included in Helion centralized logging.
 - Logging levels for vCenter should be set appropriately to prevent logging of the password for the Helion message queue.
 - The vCenter cluster and ESX Compute nodes must be appropriately backed up.
 - Backup procedures for vCenter should ensure that the file containing the Helion configuration as part of Nova and Cinder volume services is backed up and the backups are protected appropriately.
 - Since the file containing the Helion message queue password could appear in the swap area of a vCenter server, appropriate controls should be applied to the vCenter cluster to prevent discovery of the password via snooping of the swap area or memory dumps

Set up the Deployer

1. Create LUN(s), if required.
2. Download the HP Linux for HP Helion OpenStack Deployer ISO from the [Helion Downloads](#) page.
3. Boot your deployer from the ISO.
4. Enter "install" to start installation.

5. Select the language.
6. Select the location.
7. Select the keyboard layout.
8. Select the primary network interface, if prompted:
 - Assign IP address, netmask
9. Create a new account:
 - Enter a username.
 - Enter a password.
 - Enter time zone, if prompted to do so.
 - Synchronize the time on all nodes manually. NTP will be installed later.

At the end of this section you should have a deployer node set up with hLinux on it.

Configure and Run the Deployer

1. On the deployer node, enter the following command to create the SSH keypair. Please do not create the SSH Keypair if it is already present.

```
ssh-keygen -t rsa
```

2. Add ~/.ssh/id_rsa.pub to ~/.ssh/authorized_keys file:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

3. Confirm that ssh localhost works without a password and that you can get from external sources, both with and without sudo.
4. Mount the install media at /media/cdrom, for example,

```
sudo mount /media/cdrom
```

5. Unpack the following tarball:

```
tar zxvf /media/cdrom/hos-2.0.0/hlm-deployer-2.0.0-20150805T115313Z.tgz
```

6. Run the following included script:

```
~/hlm-deployer/hlm-init-2.0.0.bash
```

At the end of this section you should have a local directory structure, as described below:

helion/	Top level directory
helion/examples/	Directory contains the config input files of the example clouds
helion/my_cloud/definition/	Directory contains the config input files
helion/my_cloud/config/	Directory contains .j2 files which are symlinks to the /hlm/ansible directory
helion/hlm/	Directory contains files used by the installer

Configure Your Environment

1. Set up your configuration files, as follows:
 - a. See the sample set of configuration files in the ~/helion/examples/one-region-poc-with-esx directory. The accompanying README.md file explains the contents of each of the configuration files.

- b.** Copy the example configuration files in the required setup directory and edit them as required:

```
cp -r ~/helion/examples/one-region-poc-with-esx/* ~/helion/my_cloud/definition/
```

The configuration files for editing are available at the following location:

```
~/helion/my_cloud/definition/data
```

- The `baremetalConfig.yml` file should specify the server information for your environment.
 - The `servers.yml` file contains the IP address information for the hardware for controller nodes.
 - The `networks.yml` file contains networking information.
 - The `control_plane.yml` file contains information about the services that will be installed.
 - In the `net_interfaces.yml` file, replace all instances of “bond-mode: 1” to “bond-mode: active-backup”.
2. To continue installation copy your cloud layout to `/home/stack/helion/my_cloud/definition`
 3. Add the cloud deployment definition to the git :

```
cd /home/stack/helion/hos/ansible;
git add -A;
git commit -m 'My config';
```

4. Prepare your environment for deployment:

```
ansible-playbook -i hosts/localhost config-processor-run.yml;
ansible-playbook -i hosts/localhost ready-deployment.yml;
cd /home/stack/scratch/ansible/next/hos/ansible;
ansible-playbook -i hosts/verb_hosts site.yml;
```

5. Check the generated host files in `~/helion/my_cloud/stage/ansible/host_vars/` to ensure the correct IPs are included.

- a.** Execute the following command to review the ansible changes:

```
git show staging-ansible
```

- b.** Use the `ansible-playbook -i hosts/localhost ready-deployment.yml` file to create a deployment directory.
- c.** Run `verb_host` command from the following directory:

```
~/scratch/ansible/next/hos/ansible
```

Deploy Cobbler

1. Run the following command:

```
export ANSIBLE_HOST_KEY_CHECKING=False
```

2. Run the following playbook:

```
cd ~/scratch/ansible/next/hos/ansible
```

```
ansible-playbook -i hosts/localhost cobbler-deploy.yml
```

Provision the Nodes

1. Run the following command, which reimages all the nodes using PXE:

```
ansible-playbook -i hosts/localhost bm-reimage.yml
```

2. Wait for the nodes to install and come back up.

Deploy the ESX Cloud Controller

1. Run the following command:

```
ansible-playbook -i hosts/verb_hosts osconfig-run.yml
```

- a. Verify the network is working correctly. Ping each IP from the `/etc/hosts` file from one of the controller nodes.
2. Modify the `./helion/hlm/ansible/hlm-deploy.yml` to comment out the line containing `horizon`.
 3. Modify the `./helion/hlm/ansible/hlm-deploy.yml` to uncomment the line containing `eon.yml`.
 4. Run the following command:

```
ansible-playbook -i hosts/verb_hosts hlm-deploy.yml -e tuning_selector=medium
```

Deploy ESX Compute and OVSvAPP

The following sections describe the procedure to install and configure ESX compute and OVSvAPP on vCenter.

- [Deploy Helion Linux Shell VM Templates](#)
- [Preparation for ESX Cloud](#)

Deploy Helion Linux Shell VM Templates

The first step in deploying the ESX compute proxy and OVSvApps is to create a VM template that will make it easier to deploy the ESX compute proxy for each Cluster and OVSvApps on each ESX server.

Perform the following steps to deploy a template:

1. Import the `hlm-shell-vm.ova` in the vCenter using the vSphere client.
2. In the vSphere Client, click **File** and then click **Deploy OVF Template**.
3. Follow the instructions in the wizard to specify the data center, cluster, and node to install. Refer to the VMware vSphere documentation as needed.

Preparation for ESX Cloud

This section describes the procedures to prepare the ESX cloud for deployment.



Note: Source `service.osrc` before you perform the following procedure.

Register a vCenter Server

vCenter provides centralized management of virtual host and virtual machines from a single console.

1. Add a vCenter using EON CLI.

```
eon vcenter-add --name <vCenter Name> --ip-address <vCenter IP address> --username <vCenter Username> --password <vCenter Password> --port <vCenter Port>
```

where:

- vCenter Name - the name of the vCenter server where the service is deployed.
- vCenter IP address - the IP address of the vCenter server where the service is deployed.

- vCenter Username - the username for the vCenter administrator.
- vCenter Password - the password for the vCenter administrator.
- vCenter Port - the vCenter server port.

Sample Output:

Property	Value
created_at	2015-08-20T12:08:09.000000
deleted	False
deleted_at	None
id	BC9DED4E-1639-481D-B190-2B54A2BF5674
ip_address	10.1.200.41
name	vc01
password	<SANITIZED>
port	443
type	vcenter
updated_at	2015-08-20T12:08:09.000000
username	administrator@vsphere.local

Register ESX Cloud Network Configuration

This involve getting a sample network information template. Fill the details of the template and use that template to register cloud network configuration for the vCenter.

1. Execute the following command to get the network information template:

```
eon get-network-info-template --filename <NETWORK_CONF_FILENAME>
```

For example:

```
eon get-network-info-template --filename net_conf.json
```

Sample file of net_conf.json is shown below:

```
{
  "network": {
    # Deployer Network details
    "deployer_network": {
      #Deployer Portgroup Name.
      "deployer_pg_name": "hlm-Deployer-PG",

      #VLAN id for Deployer Portgroup
      "deployer_vlan": "1702",

      #Enable DHCP for Deployer N/W
      "enable_deployer_dhcp": "no",

      #CIDR and gateway for deployer network
      "deployer_cidr": "172.170.2.0/24",
      "deployer_gateway_ip": "172.170.2.1",
      "deployer_node_ip": "172.170.1.10"
    },

    #Management Network details
    "management_network": {
      #Mgmt DVS name.
      "mgmt_dvs_name": "hlm-Mgmt",

      #Physical NIC name for Mgmt DVS
    }
  }
}
```

```

    "mgmt_nic_name": "vmnic1",

    #Mgmt Portgroup Name.
    "mgmt_pg_name": "hlm-Mgmt-PG",

    #Interface order: Example eth1
    "mgmt_interface_order": "eth1"
  },

  "data_network": {
    #Tenant network type
    "tenant_network_type": "vlan",

    "data_dvs_name": "hlm-Data",

    "data_nic_name": "vmnic2",

    #Data Portgroup Name.
    "data_pg_name": "hlm-Data-PG",

    #Interface order: Example eth2
    "data_interface_order": "eth2",

    #If more than one mgmt_nic_name are specified then NIC teaming
    will be enabled by default
    #Active uplink NICs for NIC teaming(If this field is empty,
    first data_nic_name will be "Active" and the rest will be in "Standby")
    "active_nics": "vmnic2",

    #Load Balancing. Please choose the corresponding number
    # 1 -> Route based on the originating virtual port
    # 2 -> Route based on IP hash
    # 3 -> Route based on source MAC hash
    # 4 -> Route based on physical NIC load
    # 5 -> Use explicit failover order
    "load_balancing": "1",

    #Network Failover Detection. Please choose the corresponding
    number
    # 1 -> Link Status
    # 2 -> Beacon Probing
    "network_failover_detection": "1",

    #Notify Switches(yes/no)
    "notify_switches": "yes"
  },

  "hpcn_trunk_network": {
    #Trunk DVS name
    "trunk_dvs_name": "hlm-Trunk",

    #Trunk Portgroup Name.
    "trunk_pg_name": "hlm-Trunk-PG",

    #Interface order: Example eth3
    "trunk_interface_order": "eth3"
  },

  #VLAN Range for Data & Trunk port group. Please provide the range
  separated by a hyphen(vlan-vlan).
  #Multiple vlan or vlan ranges has to be a comma separated
  value(*OPTIONAL)
  "vlan_range": "1-4094"
},

```



```

"template": {
  #Provide the template/appliance name that will be used for cloning
  "template_name": "hlm-shell-vm"
},

"vmconfig": {
  #Number of CPUs for OVSvApp/Computeproxy VM
  "cpu": "4",

  #Amount of RAM in MB
  "memory_in_mb": "4096",

  #SSH public key content for OVSvAPP/Computeproxy password less
  login.
  "ssh_key": "ssh-rsa
AAAAB3NzaC1yc2EAAAADAQABAAQCbSGs9OeJofAp7oHrztyAWX5LKK8ZSyLjRbmPwDls0qu
+obWxRi7vJF9SdRgOB44zoLyRT2i5DC9Vz3sg4zshygLdg9qwtYTKS5N0Qi
+R8D5rnbCAPGiU7eTu3jpgVy/
xJOuCo6u1TQm2zA8epsSisqjtg6o36gZNvVcOE8XBYr92Dc3wFncjCh+Ej
+X2WsKQHiais2fgCME1g4bj2r2E4+8oTiL/
g5bhrhl1fSwQZPAMc2W018Eyum3ItHpD9stxr3OgEpR0sqk2piUasgT5lc4x9NGqa0RZgtbrEVjATBCdF6EO
qWzx9SDeD4dyDa4Y+P+Letp stack@hlm"
},
  # Do you want to skip inactive or maintenance mode hosts ?
  "skip_inactive_hosts": "yes",
  # Provide new host mo ids when you are adding a new host in an
  activated cluster.
  "new_host_mo_ids": ""
}

```

2. Modify the template (json file) as per your environment.

```
vi <NETWORK_CONF_FILENAME>
```

For example:

```
vi net_conf.json
```

3. Use the template to register Cloud Network Configuration. This sets the network information for a vCenter which is used to deploy and configure compute proxy and OVSvAPP VMs during the cluster activation.

```
eon set-network-info --vcenter-id <vCenter ID> --config-json
<NETWORK_CONF_FILENAME>
```

For example:

```
eon set-network-info --vcenter-id <vCenter ID> --config-json net_conf.json
```



Note: The vcenter ID is generated when you execute the above (step 2) command.

4. Execute the following command to view the list of clusters for the given vCenter.

```
eon cluster-list --vcenter-id <vCenter ID>
```

Sample Output

```

+-----+-----+-----+-----+
| MOID      | Name      | Datacenter | Import Status |
+-----+-----+-----+-----+
| domain-c21 | Cluster1  | DC1        | not_imported  |

```

```
+-----+-----+-----+-----+
```

Import Cluster

You can use one or more ESX clusters for ESX Cloud Deployment. When a Import Cluster is invoked, required ESX Compute Proxy and OVSvApp nodes are deployed.

1. Import the cluster for the EON database under the given vCenter.

```
eon cluster-import --vcenter-id <vCenter ID> --cluster-name <Cluster Name>
--cluster-moid <Cluster Moid>
```

where:

- vCenter ID - ID of the vcenter containing the cluster.
- Cluster Name - the name of the cluster that needs to be imported.
- cluster Moid - Moid of the cluster that needs to be imported.

Sample Output

```
+-----+-----+
| Property      | Value      |
+-----+-----+
| cpu_free       | 83071.73   |
| cpu_total      | 83072      |
| cpu_used       | 0.27       |
| datacenter     | DC1        |
| disk_free      | 1022.79    |
| disk_total     | 1023.75    |
| errors         | []         |
| memory_free    | 496.82     |
| memory_total   | 511.76     |
| memory_used    | 14.94      |
| name           | Cluster1   |
| state          | importing  |
| switches       | []         |
+-----+-----+
```

One vCenter can have multiple clusters. But it allows you to import only one cluster at a time.

Activate Clusters

When you execute the active cluster command, the `server.yml` of the input model is updated with IP Addresses of compute proxy and OVSvApp.

1. Activate the cluster for the selected vCenter.

```
eon cluster-activate --vcenter-id <vCenter ID> --cluster-moid <Cluster
Moid>
```

Sample Output

```
+-----+
+-----+
+
| Property      | Value      |
```

```

+-----+
+-----+
+
| node_info      | {u'computeproxy': {u'pxe-mac-addr':
u'00:50:56:b6:ce:1b', u'pxe-ip-addr': u'172.170.2.4', u'name':
u'COMPUTEPROXY_Cluster1', u'cluster-moid': u'domain-c21'},
u'network_driver': {u'cluster_dvs_mapping': u'DC1/host/Cluster1:hlm-
Trunk', u'Cluster1': [{u'host-moid': u'host-29', u'pxe-ip-addr':
u'172.170.2.3', u'esx_hostname': u'10.1.200.33', u'ovsvapp_node':
u'ovsvapp-10-1-200-33', u'pxe-mac-addr': u'00:50:56:b6:5e:9a'}, {u'host-
moid': u'host-25', u'pxe-ip-addr': u'172.170.2.2', u'esx_hostname':
u'10.1.200.66', u'ovsvapp_node': u'ovsvapp-10-1-200-66', u'pxe-mac-addr':
u'00:50:56:b6:56:e6'}]}} |
| resource_moid | domain-c21

```

```

+-----+
+-----+
+
| resource_name | Cluster1

```

```

+-----+
+-----+
+
| state         | activated

```

```

+-----+
+-----+
+

```

2. Execute the following command to view the list of clusters for the given vCenter.

```
eon cluster-list --vcenter-id <vCenter ID>
```

Sample Output

```

+-----+-----+-----+-----+
| MOID      | Name      | Datacenter | Import Status |
+-----+-----+-----+-----+
| domain-c22 | Cluster2 | DC1         | imported      |
+-----+-----+-----+-----+

```

Modify the Volume Configuration File

Once the cluster is activated you must configure the volume.

Perform the following steps to modify the volume configuration files:

1. Execute the following command:

```
cd /home/stack/helion/hos/ansible/roles/_CND-CMN/templates
```

2. Modify the `cinder.conf.j2` as follows:

```
# Start of section for VMDK block storage
#
# If you have configured VMDK Block storage for cinder you must
# uncomment this section, and replace all strings in angle brackets
# with the correct values for vCenter you have configured. You
# must also add the section name to the list of values in the
# 'enabled_backends' variable above. You must provide unique section
# each time you configure a new backend.
#
#[<unique-section-name>]
#vmware_api_retry_count = 10
#vmware_tmp_dir = /tmp
#vmware_image_transfer_timeout_secs = 7200
#vmware_task_poll_interval = 0.5
#vmware_max_objects_retrieval = 100
#vmware_volume_folder = cinder-volumes
#volume_driver = cinder.volume.drivers.vmware.vmdk.V
MwareVcVmdkDriver
#vmware_host_ip = <ip_address_of_vcenter>
#vmware_host_username = <vcenter_username>
#vmware_host_password = <password>
#
#volume_backend_name: <vmdk-backend-name>
#
# End of section for VMDK block storage
```

Modify the Neutron Service Configuration File

Perform the following steps to modify the neutron service configuration files:

1. Execute the following command:

```
cd /home/stack/helion/hos/ansible/roles/neutron-common/templates
```

2. Modify the `neutron.conf.j2` with the following values:

```
router_distributed = False
modify ml2_conf.ini.j2
with the below values
[ml2]
mechanism_drivers = ovsvapp, openvswitch, l2population
[agent]
enable_distributed_routing = False
```

Commit the Changes and Prepare your Environment for Deployment

1. Add the cloud deployment definition to the git :

```
cd /home/stack/helion/hos/ansible;
git add -A;
git commit -m 'My config';
```

2. Prepare your environment for deployment:

```
ansible-playbook -i hosts/localhost config-processor-run.yml;
ansible-playbook -i hosts/localhost ready-deployment.yml;
cd /home/stack/scratch/ansible/next/hos/ansible;
```

```
ansible-playbook -i hosts/verb_hosts site.yml;
```

Run the Config Processor

Execute the following commands:

```
cd ~/helion/hos/ansible  
ansible-playbook -i hosts/localhost config-processor-run.yml
```

Initiate Deploy ESX Compute Proxy and OVSvApps

Execute the following command to deploy a cloud:

```
ansible-playbook -i hosts/verb_hosts hlm-deploy.yml -e  
tuning_selector=medium
```